

Modelling Generalized Stochastic Process with Deep Generative Neural Network*

Rishav Bhagat, Erik Scarlatescu

Abstract. Modeling stochastic processes often relies on parametric methods (Wiener processes, Gaussian processes, ARCH/GARCH), where distributions are assumed known. Mixture models extend flexibility, but remain limited. Deep learning (RNNs, LSTMs, Transformers) has excelled in time series prediction, yet typically produces deterministic outputs. To capture uncertainty, deep generative models (GANs, VAEs, diffusion models) prove powerful. Non-parametric approaches (KDE, Time Series GANs) are essential when the distribution shape is unknown. Our proposed approach leverages deep learning to quantify uncertainty in stochastic process modeling, making no assumptions about the underlying distribution's form. We do this by modularizing the method into three parts: the network architecture, network stochasticity, and training scheme. We plan to run ablations and combine methods to generate better models for stochastic processes. We plan to test this approach on data from financial, physical, and biological domains, as well as on synthetic data.

Key words. stochastic processes, generative modeling

MSC codes. 68Q25, 68R10, 68U05

1. Related Work. Modeling stochastic processes has a lot of prior work, with parametric approaches like Wiener processes, Gaussian processes, and ARCH/GARCH models forming the foundation for time series analysis with their reliance on maximum likelihood estimation. These approaches provide efficient frameworks when the underlying distribution shape is known. Mixture of processes extends this concept by combining simpler experts with adaptable weighting, offering a more nuanced modeling capability [4]. However, while such parametric methods are able to represent a large process space than any of the individual experts, they are inherently limited in their ability to capture entirely arbitrary distributions.

In recent years, deep learning has revolutionized time series prediction. Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) networks were some of the first deep learning models to tackle sequential data, demonstrating their ability to capture long-range dependencies within time series. Transformer-based models, with their powerful attention mechanisms, have also gained significant traction, offering efficient learning of these dependencies and potentially surpassing RNN/LSTM performance in specific tasks. However, these models (used in their regular form), do not generate stochastic samples; they deterministically act on an input and generate a single output. [6]

Beyond deterministic prediction, deep generative models provide a compelling approach for modeling uncertainty in stochastic processes. Generative Adversarial Networks (GANs) utilize an adversarial training scheme where a generator learns to produce samples resembling the target distribution, while a discriminator attempts to distinguish real data from generated samples. Variational Autoencoders (VAEs) employ an encoder-decoder architecture with an interesting twist: noise is introduced into the latent representation during encoding, allowing the decoder to reconstruct the original data and, more importantly, generate novel samples

*Submitted to the editors DATE.

Funding: This work was funded by the Fog Research Institute under contract no. FRI-454.

that reflect the inherent uncertainty of the data distribution. Diffusion models take a unique approach, learning the process of progressively adding noise to a data point until it resembles pure noise. By reversing this learned process, they can effectively generate high-quality samples from the target distribution. [1]

Non-parametric approaches to modelling stochastic processes become necessary when the distribution shape is unknown. Kernel Density Estimation (KDE) offers another avenue for non-parametric modeling by directly estimating the distribution of future data points based on historical observations. However, KDE can encounter scalability challenges with high-dimensional data. Time series Generative Adversarial Networks (GANs) have emerged as a powerful tool, leveraging a generator-discriminator architecture in latent space to model complex stochastic processes. Here, a Wiener process injects stochasticity into the generator, allowing it to transform the input into a sample from the target process. There are many approaches applying GANs (with and without transformers) to generate time series samples. [11] [7]

This established body of prior work across various methodologies lays the groundwork for our proposed approach, which utilizes deep learning to quantify uncertainty within the context of modeling stochastic processes without any assumptions of its form.

2. Methods. Our method will consist of three modules, each with interchangeable options. We plan to run an ablation on the different combinations.

1. A flexible network architecture that supports stochastic outputs and is flexible enough to capture the complexity of arbitrary time series data
2. A method to introduce stochasticity into the network outputs to allow for sampling of stochastic processes
3. A training scheme to match the converged model’s output distribution of time series roll-outs with the data distribution.

2.1. Network Architecture. Neural networks are known to be universal function approximators, making them a suitable choice for this situation. However, this flexibility is not always an advantage, as it can mean that considerably more data is needed for neural networks to learn the patterns in the data. There exist several structures, such as RNNs, CNNs, and transformers, which build off of neural networks, while adding additional structure that serves to work as an inductive bias.

The kernels used by CNNs leverage the fact that order matters in time series data. Additionally, RNNs and transformers are well suited to sequential data. While RNNs have been largely replaced by transformers in many domains, it is still worth experimenting the former to see how well it performs in this domain.

2.2. Network Stochasticity. There are multiple ways to introduce stochasticity into network outputs. A few that we are considering right now are:

1. Concatenate a random vector $\vec{x}_r \in \mathbb{R}^k$ of length k sampled from a chosen source distribution.
2. Introduce additive noise sampled from a chosen distribution in an intermediate latent layer.

These also allow for flexibility in the chosen noise distribution. The main things we would

like to try is Gaussian noise, uniform noise, and Wiener process noise (noise conditioned on the noise from the last time step).

2.3. Training Scheme. The training schemes that could be used to learn the target distribution are variants of the GAN training scheme, variants of the VAE embedding learning, and variants of the diffusion model training scheme. We plan to modify these methods to make them fit our problem of stochastic time series sampling.

To adapt GANs to time series prediction we plan to look at methods from prior work (see section 1) for applying GANs to time series data. A simplistic method to adapt these networks to time series data is to just stack the data from the last T steps to sample the next N steps.

Each of these sections are orthogonal from each other allowing us to run ablations with many combinations and evaluate them with metrics designed to see how well the output distribution models the target distribution.

3. Applications.

3.1. Domains. The ability to model stochastic processes using neural networks has far-reaching implications in finance, physics, and biology. In finance, these models fuel more accurate stock price predictions, facilitate advanced option pricing techniques, and provide in-depth risk assessments. Within physics, neural network-based stochastic models are used to simulate complex diffusion processes, track particle movement in unpredictable environments, and explore the probabilistic realm of quantum mechanics. Biologists utilize these techniques to model population growth dynamics, analyze the intricate spread of diseases, and capture the inherent randomness within genetic processes.

Many datasets exist from these domains which can be used to test our method. However, we first plan to generate synthetic data from common stochastic processes (Wiener, Gaussian, Poisson) before trying them. We hypothesize that financial data will be the most difficult to model, due to its inherent unpredictability.

3.2. Model Applications. The model can be applied using Monte Carlo sampling methods to generate estimates of the certain statistics such as future mean and volatility/standard deviation. Additionally, traditional models that take a time series roll-out as input (such as a trading model on a financial dataset) can use samples from our time series generator to analyze properties about its distribution. Thus, using these samples, we can quantify uncertainty on these estimates by calculating variances on the metrics.

An example of such a use case is calculating returns on a trading model and then calculating uncertainty on such returns to manage risk. This can bring better estimates of Sharpe ratio and generate safer returns on average.

[1] [2] [11] [3] [4] [6] [7] [8] [12] [5] [9] [10] [13]

REFERENCES

- [1] S. BOND-TAYLOR, A. LEACH, Y. LONG, AND C. G. WILLCOCKS, *Deep generative modelling: A comparative review of vaes, gans, normalizing flows, energy-based and autoregressive models*, IEEE Transactions on Pattern Analysis and Machine Intelligence, 44 (2022), p. 7327–7347, <https://doi.org/10.1109/tpami.2021.3116668>, <http://dx.doi.org/10.1109/TPAMI.2021.3116668>.
- [2] E. BROPHY, Z. WANG, Q. SHE, AND T. WARD, *Generative adversarial networks in time series: A survey and taxonomy*, 2021, <https://arxiv.org/abs/2107.11098>.
- [3] A. DESAI, C. FREEMAN, Z. WANG, AND I. BEAVER, *TimeVAE: A variational auto-encoder for multivariate time series generation*, 2022, <https://openreview.net/forum?id=VDdDvnwFoyM>.
- [4] G. HUERTA, W. JIANG, AND M. A. TANNER, *Time series modeling via hierarchical mixtures*, Statistica Sinica, 13 (2003), pp. 1097–1118, <http://www.jstor.org/stable/24307162> (accessed 2024-03-06).
- [5] G. HUERTA, W. JIANG, AND M. A. TANNER, *Time series modeling via hierarchical mixtures*, Statistica Sinica, 13 (2003), pp. 1097–1118, <http://www.jstor.org/stable/24307162> (accessed 2024-03-06).
- [6] G. LAI, W.-C. CHANG, Y. YANG, AND H. LIU, *Modeling long- and short-term temporal patterns with deep neural networks*, 2018, <https://arxiv.org/abs/1703.07015>.
- [7] X. LI, V. METSIS, H. WANG, AND A. H. H. NGU, *Tts-gan: A transformer-based time-series generative adversarial network*, 2022, <https://arxiv.org/abs/2202.02691>.
- [8] L. LIN, Z. LI, R. LI, X. LI, AND J. GAO, *Diffusion models for time series applications: A survey*, 2023, <https://arxiv.org/abs/2305.00624>.
- [9] D. SALINAS, V. FLUNKERT, AND J. GASTHAUS, *Deepar: Probabilistic forecasting with autoregressive recurrent networks*, 2019, <https://arxiv.org/abs/1704.04110>.
- [10] K. E. SMITH AND A. O. SMITH, *Conditional gan for timeseries generation*, 2020, <https://arxiv.org/abs/2006.16477>.
- [11] J. YOON, D. JARRETT, AND M. VAN DER SCHAAR, *Time-series generative adversarial networks*, in Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver, BC, Canada, H. M. Wallach, H. Larochelle, A. Beygelzimer, F. d’Alché-Buc, E. B. Fox, and R. Garnett, eds., 2019, pp. 5509–5519, <https://proceedings.neurips.cc/paper/2019/hash/c9efe5f26cd17ba6216bbe2a7d26d490-Abstract.html>.
- [12] J. YOON, D. JARRETT, AND M. VAN DER SCHAAR, *Time-series generative adversarial networks*, in Advances in Neural Information Processing Systems, H. Wallach, H. Larochelle, A. Beygelzimer, F. d’Alché-Buc, E. Fox, and R. Garnett, eds., vol. 32, Curran Associates, Inc., 2019, https://proceedings.neurips.cc/paper_files/paper/2019/file/c9efe5f26cd17ba6216bbe2a7d26d490-Paper.pdf.
- [13] H. ZHOU, S. ZHANG, J. PENG, S. ZHANG, J. LI, H. XIONG, AND W. ZHANG, *Informer: Beyond efficient transformer for long sequence time-series forecasting*, 2021, <https://arxiv.org/abs/2012.07436>.